

FRAME-VM SPECIFICATION

Chiel Bruin

last updated: May 3, 2019

This document contains an overview of all the instructions implemented by the frame-vm and the VM itself. As the VM has two possible modes, these are discussed separately after the general machine overview.

Frame VM

As the VM is built on the concept of scopes as frames, its memory layout is made from data frames representing scopes. These data frames contain indexed slots where data can be stored. These slots correspond to declarations in the scope graph. Like scopes in a scopegraph are linked, data frames in the VM are also linked to form a graph (which can be seen as a form of a heap). Any sub-graph in this graph that is not referenced by any other part of the graph can be garbage collected, similar to garbage-collecting of values on the heap.

Besides an alternative for the heap, the VM also has a different view on the control stack. Normally this stack stores frames containing a return address, local variables and the program counter (PC). The VM is similar in that control frames are stored that contain a return address and the program counter. However, there are two big differences: Multiple return addresses are allowed (making the control-stack a control-graph) and local variables are stored in a linked data frame. In addition, there is some extra local memory that, depending on the mode, is a stack or a set of registers.

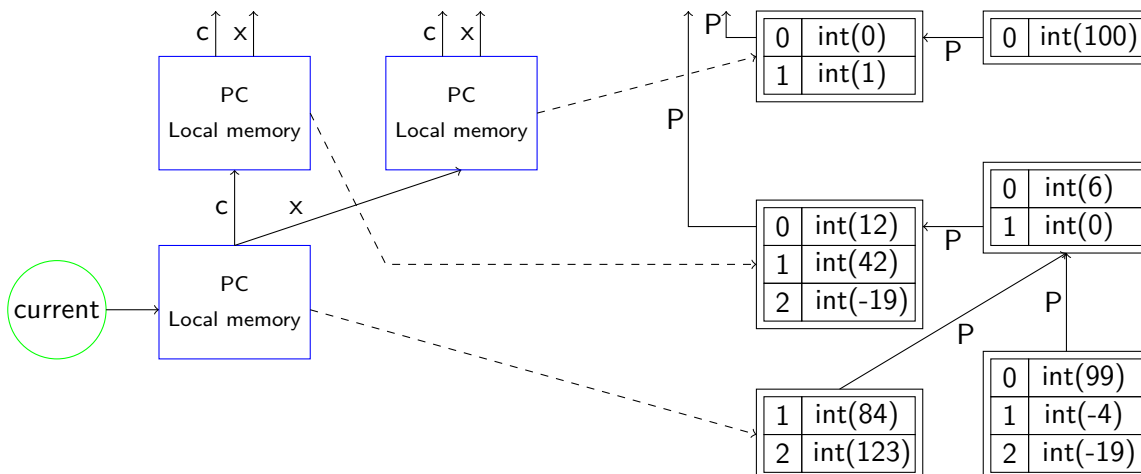


Figure 0.1: Frame VM machine layout. Control frames are displayed in blue, data frames in black. Current points to the currently executing control frame.

Having multiple return addresses for a given control frame allows you to model control-flow more easily. For example, adding exceptions is as simple as adding a second return address to the nearest exception handler. When changing to a new control frame, this extra return address needs to be

copied over (to remain pointing to the nearest handler). This process of passing around the return addresses is in a sense similar to writing execution semantics in continuation passing style (CPS).

Figure 0.1 shows a more graphical example of the layout. In this figure the machine is executing a try-catch like part of a program, as a return address (c) and nearest exception handler (x) are used.

Data Types

Possible data types that can exist in the VM, or are used in this document are:

- string: A string (only used as sugar)
- val: A generic value, can be any of the datatypes listed below
- int: An integer value
- bool¹: A boolean value
- char²: A character
- frame: A reference to a data frame
- cont: A reference to a control frame, represents an execution point (continuation)
- clos: A reference to a data frame and code block, represents a closure

¹Type alias for int

²Type alias for int

Stacy

The first mode of operation of the frameVM is stack-based. The bytecode language used in this mode is called Stacy (**stack**) and has the extension `.stc`.

For each instruction its effects on the stack are listed, together with a textual description and required arguments. After this, sugared instructions and their desugarings are listed. Understanding these reductions could provide usefull insights in the workings of the VM, but is not neccesary (assuming your language only uses function returns and exception handlers).

As the frame VM uses indexed links and slots internally, you need to define a mapping between names and indices of edge labels and continuation slots. Stacy already predefines a number of these mappings for free (namely $P \rightarrow 0$, $I \rightarrow 1$, $c \rightarrow 0$, $x \rightarrow 1$ and $n \rightarrow 2$). Adding additional labels should be done with caution.

Block syntax

In Stacy all instructions are grouped in code blocks. These blocks start with label with their unique name, followed by an indented list of instructions. This list of instructions must be followed by a control-influencing instruction to complete a block (these cannot be used inside a block). The instructions that are in this group are listed in figure 0.2

Instruction	Arguments	Instruction	Arguments
exitscope	[path] label	tailcall	label
newscope	label1 label2	tailcall	
jumpz	label1 label2	return	
jump	label	return	int
call	label1 label2	ccall	label
call	label	cret	
yield	label		

Figure 0.2: All the control-influencing instructions of Stacy. Note that some instructions in this list have a similar instruction that does not influence control.

Instructions

Table 0.1: Arithmetic operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
ipush	int		int	1	<i>Pushes the given int on the stack</i>
addi		int1, int2	int	-1	<i>Adds the two values</i>
subi		int1, int2	int	-1	<i>Subtracts int1 from int2</i>
muli		int1, int2	int	-1	<i>Multiplies the two values</i>
divi		int1, int2	int	-1	<i>Divides int2 by int1</i>
modi		int1, int2	int	-1	<i>Calculates int2 modulo int1</i>
eqi		int1, int2	bool	-1	<i>Checks if the two values are equal</i>
lti		int1, int2	bool	-1	<i>Checks if int2 is less than int1</i>
gti		int1, int2	bool	-1	<i>Checks if int2 is greater than int1</i>
ori		int1, int2	bool	-1	<i>Calculates the binary or</i>
xori		int1, int2	bool	-1	<i>Calculates the binary xor</i>
andi		int1, int2	bool	-1	<i>Calculates the binary and</i>

Table 0.2: Miscellaneous operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
print		val		-1	<i>Prints val to the console</i>
printc		frame		-1	<i>Prints frame to the console, as if it were a character</i>
prints		frame		-1	<i>Prints frame to the console, as if it were a string</i>
debug				0	<i>Generates a DOT representation of the machine state</i>
debug!				0	<i>Generates a DOT representation of the machine state and kill execution</i>

Table 0.3: Closure operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
newc	policy, lbl	frame	clos	0	<i>Creates a closure of frame with lbl as label</i>
newc	lbl	frame	clos	0	<i>Creates a closure of frame with lbl as label</i>
cnew	int	clos	cont	0	<i>Creates cont from clos with int continuation slots</i>
unpack		clos	frame	0	<i>Unpacks frame from closure</i>

Table 0.4: Type operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
int?		val	bool	0	<i>Checks if val is an integer</i>
cont?		val	bool	0	<i>Checks if val is a continuation</i>
frame?		val	bool	0	<i>Checks if val is a frame</i>
closure?		val	bool	0	<i>Checks if val is a closure</i>

Table 0.5: Frame operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
new			frame	1	Create a new frame with size 0 and pushes a reference to it on the stack
new	int		frame	1	Create a new frame with size int and pushes a reference to it on the stack
newr		int	frame	1	Create a new frame with size int and pushes a reference to it on the stack
link	[path] label	frame		-1	Link the frame on top of the stack to the given location using label as label
linkr	label	frame1, frame2		-2	Link frame2 to frame1 using label as label
copy			frame	1	Makes a shallow copy of the current frame
copy	policy1, policy2		cont	1	Makes a copy of the current execution context using policy1 for the control frames and policy2 for the data frames
copyr		frame1	frame2	0	Makes a shallow copy of frame
copyr	policy1, policy2	cont1	cont2	0	Makes a copy of cont1 using policy1 for the control frames and policy2 for the data frames
size		frame	int	0	Gets the number of slots of frame
set		val, int		-2	Store val in slot int of the current frame
set	[path]	val		-1	Store val at the given location
setr		val, int, frame		-3	Store val in slot int of frame
setr	[path]	val, frame		-2	Store val at the given location, starting path at frame
get		int	val	0	Get the value in slot int of the current frame
get	[path]		val	1	Get the value at the given location and store it on the stack
getr		int, frame	val	-1	Get the value in slot int of frame
getr	[path]	frame	val	0	Get the value at the given location, starting from frame and store it on the stack

Table 0.6: Scoping/dataframe operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
exitscope	[path]			0	Change the current dataframe to the frame at path. Breaks from nested scopes to the nesting scope
exitscope	[path] label			0	Change the current dataframe to the frame at path. Breaks from nested scopes to the nesting scope. Jump execution to label
newscope	label	frame		-1	Enters a nested scope by setting the current dataframe to frame . This new frame will be linked using label to the original frame
newscope	label1 label2	frame		-1	Enters a nested scope by setting the current dataframe to frame . This new frame will be linked using label1 to the original frame. Jumps execution to label2
mkcurrent		frame		-1	Make frame the current dataframe

Table 0.7: Control operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
jumpz	label1 label2	bool		-1	Jump to label1 if bool is false, otherwise jump to label2
jump	label			0	Unconditional jump to label
call	label1 label2	frame		0	Calls a function at location label1 using frame as execution frame. When the function returns, execution is resumed at label2
call	label	cont		0	Calls cont . When the function returns, execution is resumed at label
tailcall	label	frame		0	Calls a function at location label using frame as execution frame. Uses tail-call optimizations
tailcall		cont		0	Calls cont . Uses tail-call optimizations
return		val		-1	Return val
return	int	val{int}		-int	Return the int values on top of the stack
yield	label	val		-1	Yield val and the current continuation. Jumps execution to label
rget			val	1	Get the retruned value after a function call returns
rget	int		val{int}	int	Get int returned values after a function call returns

Table 0.8: Stack manipulation operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
pop		val		-1	Discards the element on top of the stack
dup		val	val, val	1	Duplicate the element on top of the stack
dup	int	val{int-1} val2	val2, val{int-1}, val2	1	Duplicate the element on the int-th position of the stack
swap		val1, val2	val1, val2	0	Swap the two top elements of the stack
swap	int	val1, val2 val{int-2}	val2, val{int-2}, val1	0	Swaps the element on top of the stack, with the one on the (int-1)-th position of the stack

Table 0.9: Character handling operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
spush	string		frame	1	Convert string to a character array and push it on the stack
cpush	char		char	1	Push char to the stack
printc		char		-1	Print char to the console

Table 0.10: Continuation operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
cget	[]		cont	1	Create a continuation of the current execution point
cnew	label int	frame	cont	0	Create a continuation of a new control frame with data frame frame, execution point label and size int
ccall	label	cont		-1	Call cont and set the current execution point to label
cret		cont		-1	Call cont. Do not set a next execution point
transfer	int	cont		-(int+1)	Transfer int elements as returned values to cont
transfer	int [path]			-(int)	Transfer int elements as returned values to the given continuation
cset		cont, int		-2	Store cont in slot int of the current controlframe
cset	[path]	cont		-1	Store cont at the given location
csetr		cont1, int, cont2		-3	Store cont1 in slot int of cont2
csetr	[path]	cont1, cont2		-2	Store cont1 in the given slot of cont2
cget		int	cont	0	Get the continuation in slot int of the current frame
cget	[path]		cont	1	Get the continuation at the given location
cgetr		int, cont1	cont2	-1	Get the continuation in slot int of cont1
cgetr	[path]	cont1	cont2	0	Get the continuation at the given location of cont1

Table 0.11: Exception handling operations implemented by the virtual machine

Instruction	Arguments	Pop	Push	Δ	Description
throw		val		-1	Throw the element on top of the stack to the current exception handler
try	label1 label2 label3	frame1, frame2		-2	Creates a try-catch block with frame2 as try-block running label1 and frame1 as catch-block running label2. The next instruction is at label3
try	label	cont1, cont2		-2	Creates a try-catch block with cont2 as try-block and cont1 as catch-block. The next instruction is at label

Equivalent Operations

dup \Rightarrow dup 1	return \Rightarrow return 1
swap \Rightarrow swap 1	cpush char \Rightarrow ipush char
new \Rightarrow new 0	

Figure 0.3: Simple equivalent operations

link path lbl \Rightarrow get path linkr lbl	set \Rightarrow get [] swap 2 swap setr
cnew lbl \Rightarrow get [] cnewr lbl	
get \Rightarrow get [] swap getr	set path \Rightarrow get path[:-1] swap setr path[-1:]
get path \Rightarrow get [] getr path	setr path \Rightarrow getr [i], $\forall i \in \text{path}[:-1]$ setr path[-1:]
getr path \Rightarrow getr [i], $\forall i \in \text{path}$	
getr [slot] \Rightarrow ipush slot getr	setr [slot] \Rightarrow ipush slot swap setr

Figure 0.4: Equivalent operations for frame-get, frame-set and linking

exitscope path \Rightarrow get path mkcurrent	exitscope path lbl \Rightarrow exitscope path jump lbl
newscope link \Rightarrow dup link [] link mkcurrent	newscope link lbl \Rightarrow newscope link jump lbl

Figure 0.5: Equivalent operations for control instructions (cont.)

call lbl	⇒	dup cget [x] csetr [x] dup cget [] csetr [c] ccall lbl	call lbl1 lbl2	⇒	cnew lbl1 2 call lbl2
tailcall	⇒	dup cget [x] csetr [x] dup cget [c] csetr [c] cret	tailcall lbl	⇒	cnew lbl 2 tailcall
			return n	⇒	transfer n [c] cget [c] cret
			yield lbl	⇒	cget [] swap transfer 2 [c] cget [c] ccall lbl
			throw	⇒	transfer 1 [x] cget [x] ccall

Figure 0.6: Equivalent operations for control instructions

try lbl1 lbl2 lbl3	⇒	cnew lbl2 3 dup cget [] csetr [n] dup cget [x] csetr [x] dup cget [c] csetr [c] swap cnew lbl1 3 dup cget [] csetr [n] dup swap 2 csetr [x] dup cget [c] csetr [c] dup cget [] csetr [n] ccall lbl3	try lbl	⇒	dup 2 swap dup cget [] csetr [n] dup cget [x] csetr [x] dup cget [c] csetr [c] csetr [x] dup cget [] csetr [n] dup cget [c] csetr [c] ccall lbl
---------------------------	---	--	----------------	---	--

Figure 0.7: Equivalent operations for control instructions (cont.)

Strings

The VM does not force a certain representation of arrays (as they can be cons-lists, NULL-terminated or keep track of their sizes). As an effect there is also no clear way to define how strings should be modeled. However, Stacy does provide some help when working with strings, albeit only for one of the representations and only for constructing strings. The `spush`-instruction creates a frame on the stack that contains the length of the string in slot 0, and the individual characters in consecutive slots.

This instruction can therefore be desugared in the following way:

```
spush string  ⇒  new length(string) + 1
                  dup
                  ipush length(string)
                  setr [0]
                  For all characters c at position n in string :
                  dup
                  cpush c
                  setr [n + 1]
```

Providing functionality for printing entire strings cannot be done in a similar way. This is because this functionality loops over the array and print the individual characters. Therefore the desugared version would result in multiple new code blocks that must be reused between multiple uses of the original instruction. This makes that it is more like a library function that must be included once. As this is currently not yet supported², the function should be added manually when needed.

²Work is currently done to be able to support this. Check out the development branch for a set of library functions

Helper functions

In order to aid code generation for Stacy, a number of Stratego helper strategies are provided.

- `stc-from-flat`: Given a list of Stacy instructions, generate a valid Stacy AST.
If you want to set the initial frame size, the first element of this list should be a string containing the size. If a label is found inside this list, a new block is started. This allows you to generate the code without explicitly creating code blocks (the `MAIN` label is placed before the first instruction in the list).
- `framevm-path-from-nab12`: Given a three-tuple (`name`, `namespace`, `property`) gives a Frame VM path which resolves to the declaration of `<namespace>{name}`. `property` refers to the property of the declaration where a slot index is stored.

Roger

The second mode of operation of the frameVM is register-based. The bytecode-language used in this mode is called Roger (**register**) and has the extension `.rgr`.

This language is currently still in its Alpha-phase (note the capital A), and therefore not ready for use. When the language reaches any level of (feature-)stability, this document will be updated. In short, Roger will have the same instructions as Stacy but without stack operations and the possibility to make expressions and use (control frame-local) variables.